

```
In [1]: import tensorflow as tf
        from tensorflow.keras.models import Sequential
        from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
        from tensorflow.keras.preprocessing.image import ImageDataGenerator
        import matplotlib.pyplot as plt
        import numpy as np
```

```
In [5]: # Paths to the directories
        train_dir = 'train'
        val_dir = 'val'
        test_dir = 'test'
```

```
In [6]: # 1. Data Preprocessing and Augmentation
        train_datagen = ImageDataGenerator(rescale=1./255, zoom_range=0.2, horizontal_flip=
        val_datagen = ImageDataGenerator(rescale=1./255)
        test_datagen = ImageDataGenerator(rescale=1./255)
```

```
In [7]: # Loading images from 'train' directory
        train_generator = train_datagen.flow_from_directory(
            train_dir,
            target_size=(150, 150), # Resize images to 150x150
            batch_size=32,
            class_mode='binary' # Binary classification (pneumonia vs normal)
        )
```

Found 5216 images belonging to 2 classes.

```
In [8]: # Loading images from 'val' directory
        val_generator = val_datagen.flow_from_directory(
            val_dir,
            target_size=(150, 150),
            batch_size=32,
            class_mode='binary'
        )
```

Found 16 images belonging to 2 classes.

```
In [9]: # Loading images from 'test' directory
        test_generator = test_datagen.flow_from_directory(
            test_dir,
            target_size=(150, 150),
            batch_size=32,
            class_mode='binary'
        )
```

Found 624 images belonging to 2 classes.

```
In [10]: # 2. Build the CNN model
        model = Sequential([
            Conv2D(32, (3, 3), activation='relu', input_shape=(150, 150, 3)),
            MaxPooling2D(pool_size=(2, 2)),

            Conv2D(64, (3, 3), activation='relu'),
            MaxPooling2D(pool_size=(2, 2)),
```

```

Conv2D(128, (3, 3), activation='relu'),
MaxPooling2D(pool_size=(2, 2)),

Flatten(),
Dense(128, activation='relu'),
Dense(1, activation='sigmoid') # Binary classification (pneumonia vs normal)
])

```

C:\Users\bonde\anaconda3\Lib\site-packages\keras\src\layers\convolutional\base\_conv.py:107: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```

super().__init__(activity_regularizer=activity_regularizer, **kwargs)

```

```

In [11]: # 3. Compile the model
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

```

```

In [12]: # 4. Train the model with validation
history = model.fit(
    train_generator,
    epochs=10,
    validation_data=val_generator # Use the validation set for evaluation during t
)

```

Epoch 1/10

C:\Users\bonde\anaconda3\Lib\site-packages\keras\src\trainers\data\_adapters\py\_dataset\_adapter.py:122: UserWarning: Your `PyDataset` class should call `super().\_\_init\_\_(\*\*kwargs)` in its constructor. `\*\*kwargs` can include `workers`, `use\_multiprocessing`, `max\_queue\_size`. Do not pass these arguments to `fit()`, as they will be ignored.

```

self._warn_if_super_not_called()

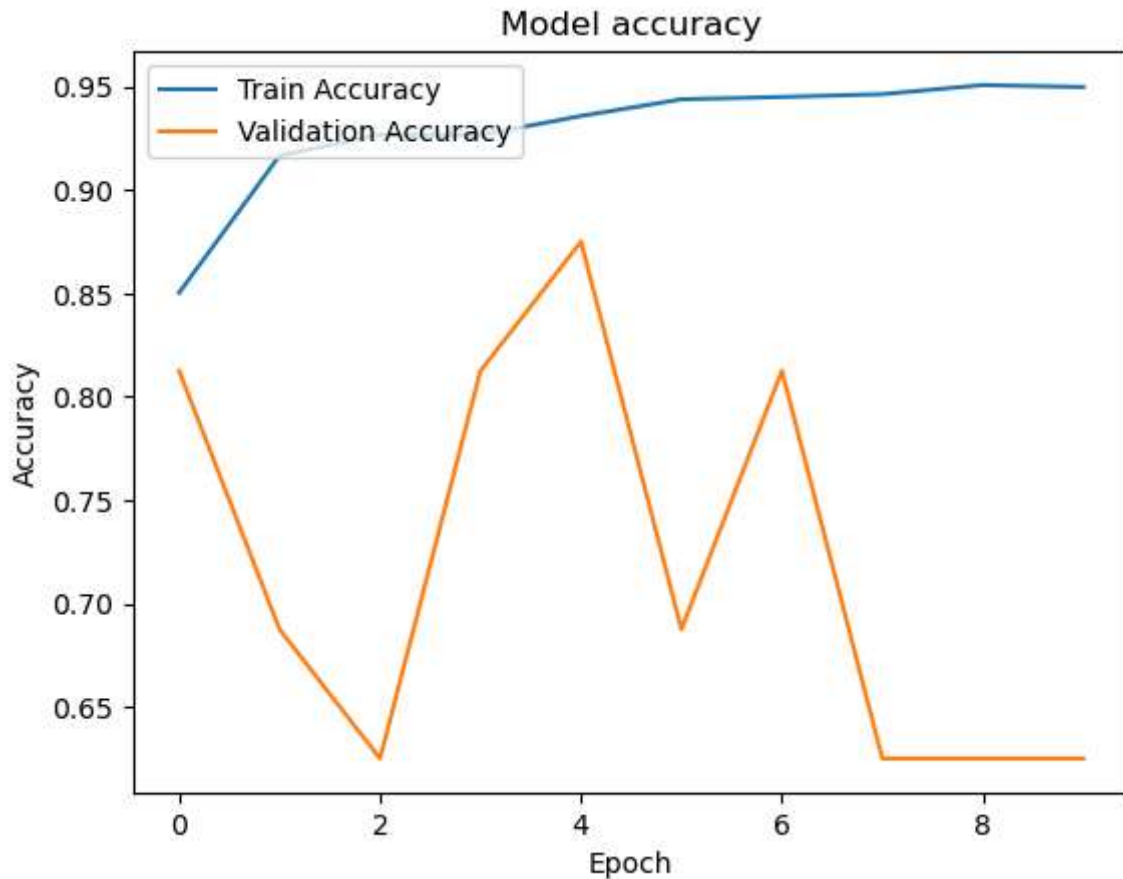
```

163/163 ————— 106s 612ms/step - accuracy: 0.7878 - loss: 0.4909 - val\_accuracy: 0.8125 - val\_loss: 0.4107  
Epoch 2/10  
163/163 ————— 97s 574ms/step - accuracy: 0.9130 - loss: 0.2123 - val\_accuracy: 0.6875 - val\_loss: 0.7915  
Epoch 3/10  
163/163 ————— 97s 574ms/step - accuracy: 0.9304 - loss: 0.1769 - val\_accuracy: 0.6250 - val\_loss: 0.8974  
Epoch 4/10  
163/163 ————— 97s 576ms/step - accuracy: 0.9242 - loss: 0.1762 - val\_accuracy: 0.8125 - val\_loss: 0.4453  
Epoch 5/10  
163/163 ————— 97s 573ms/step - accuracy: 0.9376 - loss: 0.1580 - val\_accuracy: 0.8750 - val\_loss: 0.5180  
Epoch 6/10  
163/163 ————— 127s 759ms/step - accuracy: 0.9427 - loss: 0.1511 - val\_accuracy: 0.6875 - val\_loss: 0.5395  
Epoch 7/10  
163/163 ————— 100s 593ms/step - accuracy: 0.9376 - loss: 0.1481 - val\_accuracy: 0.8125 - val\_loss: 0.2721  
Epoch 8/10  
163/163 ————— 103s 613ms/step - accuracy: 0.9444 - loss: 0.1431 - val\_accuracy: 0.6250 - val\_loss: 1.5458  
Epoch 9/10  
163/163 ————— 108s 646ms/step - accuracy: 0.9512 - loss: 0.1316 - val\_accuracy: 0.6250 - val\_loss: 1.1715  
Epoch 10/10  
163/163 ————— 127s 751ms/step - accuracy: 0.9491 - loss: 0.1308 - val\_accuracy: 0.6250 - val\_loss: 0.9418

```
In [13]: # 5. Evaluate the model on the test set
test_loss, test_acc = model.evaluate(test_generator)
print(f"Test Accuracy: {test_acc:.3f}")

# 6. Plot training & validation accuracy
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(loc='upper left')
plt.show()
```

20/20 ————— 8s 390ms/step - accuracy: 0.8440 - loss: 0.5190  
Test Accuracy: 0.864



```
In [14]: # Make predictions on the test data
predictions = model.predict(test_generator)

# Get the class labels for binary classification (NORMAL = 0, PNEUMONIA = 1)
class_labels = list(test_generator.class_indices.keys())
```

20/20 ————— 7s 332ms/step

```
In [17]: # Function to plot images with their predictions
def show_predictions(generator, num_images=10):
    plt.figure(figsize=(15, 15))

    for i in range(num_images):
        # Get a batch of test images and Labels
        img_batch, label_batch = next(generator) # Use next() to get images and La
        img = img_batch[0] # Get the first image in the batch
        label = label_batch[0] # Get the first label in the batch


        # Predict on the single image
        prediction = model.predict(np.expand_dims(img, axis=0)) # Make prediction
        predicted_class = (prediction > 0.5).astype("int32")[0][0]

        # Plot the image and show prediction
        plt.subplot(5, 5, i + 1)
        plt.imshow(img)
        plt.title(f'Pred: {class_labels[predicted_class]}\nTrue: {class_labels[int(
        plt.axis('off')
```

```
plt.show()
```

```
In [18]: # Show 10 predictions
show_predictions(test_generator, num_images=10)
```

```
1/1 ————— 0s 127ms/step
1/1 ————— 0s 34ms/step
1/1 ————— 0s 31ms/step
1/1 ————— 0s 18ms/step
1/1 ————— 0s 18ms/step
1/1 ————— 0s 15ms/step
1/1 ————— 0s 16ms/step
1/1 ————— 0s 47ms/step
1/1 ————— 0s 16ms/step
1/1 ————— 0s 16ms/step
```



```
In [ ]:
```